



**ΠΑΝΕΠΙΣΤΗΜΙΟ
ΙΩΑΝΝΙΝΩΝ**

Enterprise Integration Patterns Building message-oriented middleware with Apache Camel

Odysseas Neslechanidis
Supervisor: Christos Gkogkos

August 15, 2022

Graduation Thesis

University of Ioannina
Department of Informatics and Telecommunications

Abstract

The term "Enterprise Integration Patterns (EIPs)" refers to a vocabulary of solutions to common problems in the integration of enterprise systems. Of such vocabularies pattern languages may be constituted to allow complex business flows of diverse form to be described and handled in a uniform way.

Apache Camel is a framework that implements EIPs around a common interface based on Java Message Objects. Camel also provides an IDE-friendly declarative Domain Specific Language (DSL) oriented around this interface, which enables integration flows between disparate systems ("Camel routes") to be described neatly as Java Messages passed around between chained camel methods.

The specifics of the underlying communication protocols (FTP, http, ActiveMessageQueue etc) are abstracted away and the flow of information is cleanly described, leaving such considerations as availability, load balancing, validation, security as the primary factors influencing the middleware's architectural complexity.

In this thesis production deployments of Java Spring middleware utilizing Apache Camel will be studied. The most commonly used EIPs' Camel implementations will be inspected, and a comparison with more established integration tooling will be made when convenient, to ascertain the benefits of the Message-Oriented Middleware (MOM)-backed Camel DSL approach.

This thesis was approved by a three-person examination committee.

Examination Committee

1. Christos Gkogkos
2. John Doe
3. John Smith

Affidavit

I hereby affirm that this Bachelor's Thesis represents my own written work and that I have used no sources and aids other than those indicated. All passages quoted from publications or paraphrased from these sources are properly cited and attributed. The thesis was not submitted in the same or in a substantially similar version, not even partially, to another examination board and was not published elsewhere.

Signed,
Neslechanidis Odysseas

All Rights Reserved ©

Contents

I	The language of Enterprise Application Integration (EAI)	5
1	Introduction	5
2	Islands of automation and the advent of EAI	5
3	Introducing EAI in a organization	7
3.1	General challenges	7
3.2	Types of integration	8
3.2.1	Information Portals	8
3.2.2	Data Replication	8
3.2.3	Shared Business Functions	9
3.2.4	Service-Oriented Architectures and Distributed Business Processes	9
3.2.5	Business-to-Business Integration	9
4	Integration in a Service-Oriented Architecture	9
5	Messaging in practice: Message Oriented Middleware	11
6	Enterprise Integration Patterns	13
6.1	Messaging	13
6.1.1	Message channel	13
6.1.2	Router	13
6.1.3	Message endpoint	13
II	Apache Camel: EIPs in action	13
7	Introduction	13
8	Terminology	13
9	Apache Camel field study	14

Part I

The language of Enterprise Application Integration (EAI)

1 Introduction

Enterprise Application Software (EAS) is the term for computer programs used to satisfy the needs of an organization rather than individual users. Almost all business operations, at different points in time, have come to benefit from the proliferation of software in this space. Commonly used acronyms used to categorize such software include ERP (Enterprise Resource Planning), CRM (Customer Relationship Management), BI (Business Intelligence), CMS (Content Management System), WMS (Warehouse Management System). They serve to automate every business need of modern enterprises, from its customer facing operations, to keeping track of warehouse inventory, calculating billing and taxes, observing regulations, and much more. While comprehensive enterprise software suites offering differing degrees of customizability have come to exist, owing to the organisational similarity of enterprises above a certain scale, switching costs (CITE), preservation of optionality in partnering with software vendors (CITE EIPbook p32), as well as other adjoining business considerations, have hindered their more widespread adoption. This has necessitated the systematic study and development of solutions for Enterprise Application Integration, among which message-oriented middleware has stood out as one the most promising.

2 Islands of automation and the advent of EAI

The term “Islands of automation” was a popular term introduced in the 1980s to describe the status quo of automation systems existing within information silos. The rapid development and adoption of enterprise software systems during this time came to pass with little regard for the ability of those systems to communicate with one another.

(TODO INNEFICIENCIES of isolation and drawbacks of adhoc point2point solutions , see archpatbook p25, 351, 252

- p2p being massively complex, only suitable for asynchronous solutions

- hubnspoke helping reduce complexity by providin a common integration platform. event-driven

- service exposure enabling synchronous communication. daemon/service listener driven (request-response)

- the two above, combined in an integration runtime, with added logic for mediation/obscuring between the uniform exposure protocols and data formats, and the legacy formats and protocols (or informal specs) of the integrated applications, constitute an ESB.

integration engines of today allow for the implementation of different patterns that are more lightweight and decentralized. no longer is runtime cloning of the entire integration engine required for scaling infrastructure, nor is there hard need for a dedicated team with deep expertise in the particular SOA solution. how this is achieved:

-
an essential addition to the ESB pattern were the service registry capabilities of modern integration runtimes, as relying on interface documentation e.g. swagger or human2human interaction erodes reusability, which was the initial promise of hubnspoke and SOA in general.

TODO where does oauth fit?

service exposure gateways sometimes take up the role of service registries and also often implement other API management capabilities, such as encryption/redaction, traffic control, a security model with patterns for access control, e.g. with user roles/permission schemes etc.

The gateways can be supplemented with web portals that describe the available APIs, enable potential users to issue keys automatically (self-subscribe) in order to use the APIs, and provision analytics for both users and providers of the APIs.

TODO “Increasingly, more modern systems of record already provided an HTTP-based interface that only needed controlled exposure using the exposure gateway. The integration runtime was only required when more complex integration took place, such as more unusual protocols, data formats, compositions of multiple requests—or perhaps in cases where transactionality was needed.

The introduction of an API management layer led to the obvious question: What now is the ESB? Many had come to see the integration runtime and the ESB pattern as one and the same. But in fact, if the ESB pattern is all about exposing services and APIs, then the boundaries of the pattern really include both the integration runtime and the exposure gateway, and in some cases just the gateway. However, due to the ESB’s incorrect association with the integration runtime, we have to accept that this is not how the ESB term is typically used.”

back-end-for-front-end (BFF) pattern (present in mobile apps/ SPAs) was the first foray outside enterprise: APIs perfectly suited to the needs of a prospective frontend, with rationalized data models, ideal granularity of operations, specialized security models etc

from then on, public exposure of partner APIs has steadily become more common, providing the opportunity for collaboration among loosely parties. The resultant market-like ecosystem has been dubbed the “API Economy” TODO speak about it, check 2 ebooks . The emergence of this developmental trend arguably does credit to the promise of reusability the idea of web service orientation has borne since it’s inception.

TODO SaaS perhaps deserves a mention here

)

The field of Enterprise Application Integration (EAI) is a field of study aiming to refine a framework for rectifying these inefficiencies. The shifting nature

of the business landscape and of enterprises that operate within it, together with the continued innovation in, and expansion of, the EAS space, has resulted in it being a complicated problem to tackle.

Enterprise software is adopted at different times, it is developed from different vendors, at different points in time, oriented towards different business needs.

The employment of Domain-driven design, in recognition of the maintainability and extensibility benefits domain-expert input in the refining of the software's domain model confers, is a fact further complicating the effort of business software consolidation.

To this day, the introduction of a (TODO define middleware) middleware stack remains a very common business need.

3 Introducing EAI in a organization

3.1 General challenges

Prior to engaging with the technical aspects of Enterprise Application Integration, it is necessary to consider a set of social and organizational challenges that the development and adoption of such a solution might necessitate or bring about.

Enterprise Application Integration often requires a significant shift in corporate politics. By extension of Conway's law that postulates that "Organizations which design systems are constrained to produce designs which are copies of the communication structures of these organizations.", it appears that the consolidation of enterprise software tools serving business processes often necessitates a consolidation of the business units and IT departments involved in those same processes. (CITE EIPbook p32)

Furthermore, owing to the wide scope of a middleware integration solution bringing together critical business functions, the novel risk of failure or misbehavior of such a system has to be internalized. The risk profile and magnitude of reorganization around such a single point of failure ought to be carefully considered.

Bordering the technical side, the feasibility of integrating systems by modifying them to better fit the integration architecture, rather than by having to design the integration architecture to work around the various systems' limitations and deficiencies, also often depends on political factors. In that vein, unsupported legacy systems still in operation, systems under proprietary licenses, and systems whose support is outsourced under more or less stringent long-term agreements can adversely influence the complexity of the final product.

In terms of standardization, it bears mentioning that despite the benefit of convergence around Web Services and a Service Oriented approach to middleware architecture (which will be expounded upon in later chapters), the proliferation of new extensions or interpretations of the standard, and most significantly the shift towards REST (and, more recently GraphQL) in lockstep with the mo-

mobile revolution, has created new challenges for integration engineers. REST, in particular, owing to it being an architectural style for software that expose http APIs rather than a protocol for web services per se, is frequently implemented partially and/or wrongly, often necessitating ad-hoc code for the consumption of APIs exposed in this manner.

Finally, the operations aspect of utilizing middleware solutions presents a unique challenge, as maintenance, deployment, monitoring and troubleshooting of such heterogeneous, distributed systems commonly require mixes of skills which are not par for the course to be found in single individuals. To companies or organizations of sufficient scale as to already necessitate a formalized employee training regime, the overhead for the maintenance of such human capital might be lower. (CITE EIPbook p32)

3.2 Types of integration

While the above challenges generally apply to every approach in the broader category of integration, many further issues have to be considered depending on the business aims that dictate, and the technical aspects that come as a consequence of, the prospective type of integration solution. The following categorization has been proposed:

3.2.1 Information Portals

Information portals serve to aggregate information from disparate systems within an organization with the aim of making it more accessible to humans. They often facilitate the collaboration between different departments and physical locations. They are also commonly used in business decision processing and data analysis. Common features include multi-window views serving information from different sources with automatic refresh of related windows during navigation, search, tagging and other categorization schemes.

Various other more advanced features are common, but being as they cater to particular business functions, employees roles or departments, no account of those will be attempted. Indeed, one of the common abstract features, or aims, of such systems, is the personalization of the displayed information, achieved through the profiling of users based on role, experience, competencies, habits and expressed preferences.

3.2.2 Data Replication

Many business systems require access to the same data, but are designed to utilize their own, separate datastores. The resulting data replication necessitates provisions for maintaining the data synchronized. Commonly utilized for those purposes are the replication features built into modern Database Management Systems, the file export and import functions supported by many Enterprise Software Systems, and message-oriented middleware automating transport of data via messages between arbitrary datastore solutions.

3.2.3 Shared Business Functions

Needless duplication can exist in code serving business functions as well as in data. Were supported, invocation of shared business functions implemented as services (A service is a well-defined function that is universally available and responds to requests from “service consumers”) can help avoid the native implementation of redundant functionality.

Were feasible, the need for data replication can also be circumvented via this approach by serving shared data as a service. In that vein, some criteria to be considered include the amount of control that is had over the systems (calling a shared function is usually more intrusive than loading data into the database) and the rate of change of the relevant data (service invocation is costlier than data access, therefore is less efficient for relatively frequently accessed, relatively static data).

3.2.4 Service-Oriented Architectures and Distributed Business Processes

Once an enterprise assembles a collection of useful services, managing the services becomes an important function.

Service Oriented Architecture is a proposed style of service design and orchestration that incorporates the best industry practices in structuring middleware solutions around services that correspond to business functions. This particular approach to middleware architecture shall be expounded upon in a later chapter.

A variant dubbed “Distributed Business Process”, is also to be found in the bibliography. It concerns the design of management services that serve to coordinate the execution of the relevant business functions that are implemented natively in an integrated system’s constituent applications, in order to achieve every particular business process. Such schemes can exist within larger SOA-abiding systems, and the lines between the two terms often blur.

3.2.5 Business-to-Business Integration

In many cases, business functions may be available from outside suppliers or business partners. Business to Business (B2B) integration software provides the architecture needed to digitize information and route it through an organization’s trading ecosystem (usually online platforms) using the Electronic Data Exchange (EDI) format appropriate for the application.

4 Integration in a Service-Oriented Architecture

Service Oriented Architecture (SOA) is an evolution of predecessors such as component-based architecture and Object Oriented Analysis and Design of remote objects e.g. the CORBA standard.

Component-based architecture, or component-based software engineering, emphasizes separation of concerns with respect to the various functions provided

in a given software system. Components are commonly implemented around interfaces, that encapsulate the particulars of the components' implementation, and narrow the available surface-area for wiring together the various functionally autonomous modules. Cohesion is maintained by fitting additional modules onto the interfaces. The modules, which can be of arbitrary origin, are rendered into components by implementing their respective interfaces. The modules can exist as components locally within the same virtual or physical machine, or in the context of distributed systems such as networks (e.g. as web services or web resources).

In the SOA evolution of this approach, reusability and use in the context of distributed systems is emphasised. To realize this architectural style's potential, the promulgation of Web Service standards becomes instrumental. In this way it is ensured that networked software components can be developed as generic "Web Services", or business function-specific components that are implemented without knowledge or regard for the multitude of systems in which they may be utilized.

The effort towards this end has borne results through the W3C Web Services specification, though nowadays the emergence of alternatives and the REST architectural style in particular has created a rift in the SOA ecosystem, which is nevertheless efficiently bridged by another core feature of SOA, the Enterprise Service Bus.

-TODO LITMUS

and the establishment of the practice of environment definition via Web Service Litmus tests. (TODO litmus is about

http://deg.egov.bg/LP/soa.rup_soma/tasks/soa_service_qualification_E0D920A6.html

constraining the potential services to a subset that is reusable in the environment in which it is initially developed),

As mentioned previously, a distinctive feature of the SOA style is business centeredness, with components or interfaces aimed at fitting business functions, rules, or goals.

Service-oriented architecture can be implemented with web services or Microservices. This is done to make the functional building-blocks accessible over standard Internet protocols that are independent of platforms and programming languages.

(TODO domain driven design and microservices after loose coupling, reusability emphasis, ESBs) . It's central idea is the creation of a Ubiquitous Language with the assistance of domain experts, that embeds domain terminology into the software components' naming and structure. This approach is suitable for the design of middleware as it helps confer flexibility and extensibility to the resulting system, and reduces friction in the operations side, which, as mentioned previously, is an important factor when considering the adoption of integration solutions.

Each SOA system defines it's own environment, by which every service implementation must abide. This is ensured through the provision of a "Litmus Test", that determines whether a given service implementation is correct in it's particular SOA system's context.

A further recommendation of the SOA style is that open standards be used, their use being instrumental in realizing interoperability with different consumer implementations, and location transparency.

The related buzzword

https://en.wikipedia.org/wiki/Service-oriented_architecture#Defining%20concepts

https://en.wikipedia.org/wiki/Enterprise_service_bus

A further property service-orientation promotes is loose coupling between services. SaaS can be considered to have evolved from SOA.

Loose coupling, in addition to enabling the development of distributed architectures composed of programs developed by different teams at different times, allows for domain-driven design to be observed in its constituent parts, which is claimed to increase maintainability and creative cross-domain collaboration.

Domain-Specific-Languages and Aspect-Oriented Programming can be used to manage the complexity produced by the increased need for isolation and encapsulation that Domain-Driven design necessitates.

https://en.wikipedia.org/wiki/Domain-driven_design

Loose coupling is achieved through transactions, queues provided by message-oriented middleware, and interoperability standards.

Transactions help ensure validity of exchanges, queues enable asynchronicity and load balancing in distributed systems, and interoperability standards provide a common target for the integration of legacy systems (often rendering them network-enabled in the process) and newly implemented services alike.

(TODO benefits of SOA https://en.wikipedia.org/wiki/Service-oriented_architecture#Organizational_benefits)

In the messaging approach, provisions for asynchronicity (message buffers, brokers) and arbitrary consumer scaling are made.

(TODO Pure Messaging Integration approach https://en.wikipedia.org/wiki/Apache_ActiveMQ#Usage)

Hybrid Web Service - Messaging SOAs using MOMs is common practice.

(TODO benefits with Microservice-based SOA https://en.wikipedia.org/wiki/Service-oriented_architecture#Implementation_approaches)

Microservices is a novel, implementation agnostic approach to SOA, that allows for domain-driven design to be observed (actually they are only loosely related and, in fact, operate at different scopes, as discussed here

<https://www.ibm.com/cloud/learn/soa>)

https://en.wikipedia.org/wiki/Service-oriented_architecture#Microservices

5 Messaging in practice: Message Oriented Middleware

(TODO read EIPbook forewords on SOA, asynchronous messaging of self reliant systems being the point of integrated systems as opposed to n-tier codependent distributed systems, conferring benefits of request throttling and load balancing

but increasing complexity; this fact makes asynchronous messaging a promising approach, and this has informed the direction of this thesis)

(TODO EIPbook p64 on Integration Approaches, with messaging as the most promising one)

(TODO SOA is enabled via messaging. Time to get deep into it)

Apache Camel is a framework for building MOM middleware. More generally, it aspires to enable integrations designed around the Enterprise Integration Pattern (EIP) vocabulary. In addition to native support for ActiveMQ and other message brokers via JMS, it provides features that enable most common SOA architectures, modern and legacy alike. Standard SOAP Web Services, RESTful http Web Services and more are natively supported, with Amazon Web Services, GraphQL and other modern technologies supported as Extensions.

Examples from Java OOP, JMS messaging, Camel.

- Common architectural elements of MOM systems (request databases/message buffers, aggregators, api consumers, services).
- Intruduction to messaging, key problems it solves (separation of concerns, decoupling etc; use analogues from different domains e.g. URI barcodes etc) Identifier vs Locator disambiguation

Message passing implies URLs (Uniform Locator of Resources) Uniform Identifier: as in URI: 23-digit barcode form unique id of a thing whose location or mode of access is not defined plus Resource Locator:

-Locator as in http:// (mode of access) which implies the rest is an address.

-Resource as in shared-nothing, volatile object: Evoke the same operation (e.g INSERT) twice, don't expect it to be the same function call because the message-passing abstraction hides (therefore isolates) underlying state changes.

Eg. a necessary rerouting of a request because a node is down does not concern the message sender. Drawing out the metaphor, URI = class interface, URL = address of volatile object

- Message buffer considerations: When full: - Block the sender (deadlock risk) - Drop future messages (producer-consumer problem; unreliability), Asynchronicity and concurrency: gotta have both! + friends
- Services as components of monoliths vs distributed systems. Messaging in OOP vs JMS topic, queue schema. Compare and contrast, justify differences by comparing problem domains.

Pros: Shared--nothing, all the loose-coupling stuff

Cons: overhead as arguments need to be copied and transmitted to the receiving object

- https://en.wikipedia.org/wiki/Message_passing#Distributed_objects

6 Enterprise Integration Patterns

6.1 Messaging

Message (and exchange), pipes and filters, router, translator, endpoint

6.1.1 Message channel

point2point, publish-subscribe, channel adapter, message bus

6.1.2 Router

filter, splitter, aggregator, broker

6.1.3 Message endpoint

consumers: polling, event-driven, message dispatcher

Part II

Apache Camel: EIPs in action

7 Introduction

Talk about Camel being a tool for working with EIPs, what it's built on, tooling eg IDEs and frameworks that support it, companies that invest on it and custom products eg JBoss

8 Terminology

A Component is a factory for creating Endpoint instances.

Processors are necessarily implementations of the Camel Processor interface. Apart from that, services differ from processors in that they are meant to be used procedurally, called directly from within other methods, whereas processors are integrated in a program's flow through message-based integration, which is what Camel is built for.

This makes services the proper abstraction for program-wide configuration injection (e.g. `PartnerManagementService`) and a feasible one for utilities for which there is no expected need for integration with remote components (e.g. `RequestHandlingService`).

- Server as persistence hub, interacting with client services via library calls with "direct:" route chaining.
- Split, choice and aggregator branching and joining.

- “from:”-driven abstract pollers with different implementations per instance.
How can Camel help adapt towards a more scalable microservice variant?

Multiple task scheduling components are available. They produce timer events that can be used to trigger recurring camel routes via consumer EIPs, or otherwise provide a means of time tracking for local or distributed tasks. The primary ones are scheduler (or it’s simpler variant, timer) and quartz.

The scheduler component utilizes the host jdk’s timer and is intended for locally tracked tasks that have no need for accuracy, as no provision is made against downtime.

The quartz component uses a database to store timer events and supports distributed timers, and is therefore fault tolerant and suitable for scheduling distributed tasks.

declarative DSLs/xml

(TODO <https://martinfowler.com/dsl.html>)

Examples from declarative vs procedural DSLs (e.g. shell vs guile scripts), Spring traditional remoting vs Camel, timer/from components introduced here, as an example of how Camel departs from older integration methods..

- Producer, consumer properties per component.
- Seamless remoting via “.to()” chaining. Direct ProducerTemplate calls.
- Declarative programming advantages, mention drag-and-drop services (e.g. redhat’s integration product).
- Machine readable markup vs Camel’s DSL hack.

9 Apache Camel field study

Examining a 3PL logistic company’s middleware stack

References

- [1] Christos Gogos, Angelos Dimitzas, Vasileios Nastos, and Christos Valouxis. Some insights about the uncapacitated examination timetabling problem. In *2021 6th South-East Europe Design Automation, Computer Engineering, Computer Networks and Social Media Conference (SEEDA-CECNSM)*, pages 1–7. IEEE, 2021.